

# Complete Guide to Dataset Preprocessing for Machine Learning

## Contents

|          |  |           |
|----------|--|-----------|
| <b>1</b> | <b>What is Data Preprocessing?</b>                 | <b>2</b>  |
| <b>2</b> | <b>Common Data Issues</b>                          | <b>2</b>  |
| <b>3</b> | <b>Preprocessing Techniques</b>                    | <b>2</b>  |
| 3.1      | Handling Missing Values                            | 2         |
| 3.1.1    | Remove Rows/Columns with Missing Values            | 2         |
| 3.1.2    | Imputation   | 2         |
| 3.1.3    | Using Algorithms to Predict Missing Values         | 3         |
| 3.2      | Data Cleaning                                      | 3         |
| 3.2.1    | Remove Duplicates                                  | 3         |
| 3.2.2    | Fix Inconsistent Formats                           | 4         |
| 3.2.3    | Handle Incorrect Data                              | 4         |
| 3.3      | Encoding Categorical Variables                     | 4         |
| 3.3.1    | Label Encoding                                     | 4         |
| 3.3.2    | One-Hot Encoding                                   | 4         |
| 3.3.3    | Target Encoding                                    | 5         |
| 3.4      | Feature Scaling                                    | 5         |
| 3.4.1    | Standardization (Z-Score Normalization)            | 5         |
| 3.4.2    | Min-Max Scaling                                    | 5         |
| 3.5      | Handling Outliers                                  | 6         |
| 3.5.1    | Remove Outliers                                    | 6         |
| 3.5.2    | Cap Outliers                                       | 6         |
| 3.5.3    | Transform Features                                 | 6         |
| 3.6      | Feature Selection                                  | 7         |
| 3.6.1    | Filter Methods                                     | 7         |
| 3.6.2    | Wrapper Methods                                    | 7         |
| 3.6.3    | Embedded Methods                                   | 7         |
| 3.7      | Feature Engineering                                | 8         |
| 3.7.1    | Creating New Features                              | 8         |
| 3.7.2    | Transforming Features                              | 8         |
| 3.8      | Data Splitting                                     | 8         |
| <b>4</b> | <b>Putting It All Together: A Complete Example</b> | <b>9</b>  |
| <b>5</b> | <b>Best Practices and Tips</b>                     | <b>10</b> |

## 1 What is Data Preprocessing?

Data preprocessing involves preparing raw data for machine learning by cleaning, transforming, and organizing it. Machine learning models rely on numerical data, so preprocessing ensures the data is consistent, complete, and scaled appropriately. Poor preprocessing can lead to inaccurate models, while good preprocessing improves model performance.

Think of preprocessing as preparing ingredients before cooking. You need to wash, chop, and measure ingredients (data) to ensure the dish (model) turns out well.

## 2 Common Data Issues

Before diving into techniques, let's understand the common problems in datasets:

- **Missing Values:** Some data points are absent (e.g., empty cells in a CSV file).
- **Inconsistent Data:** Typos, mixed formats (e.g., "USA" vs. "United States"), or duplicates.
- **Categorical Data:** Non-numeric data (e.g., "Male"/"Female") that models can't directly process.
- **Outliers:** Extreme values that don't align with the rest of the data.
- **Unscaled Features:** Features with different ranges (e.g., age: 0-100, salary: 0-100000) can bias models.
- **Irrelevant Features:** Columns that don't contribute to the prediction.
- **Imbalanced Data:** When one class dominates (e.g., 90% "No" vs. 10% "Yes" in a dataset).

Preprocessing addresses these issues to make the data model-ready.

## 3 Preprocessing Techniques

### 3.1 Handling Missing Values

Missing values (e.g., NaN in Python) can cause errors in machine learning models. Here are common strategies to handle them:

#### 3.1.1 Remove Rows/Columns with Missing Values

- **When to Use:** If only a small percentage of data is missing, or if a column has too many missing values to be useful.
- **Pros:** Simple and avoids introducing bias.
- **Cons:** Can lose valuable data if too many rows are removed.

#### 3.1.2 Imputation

Replace missing values with a statistic like mean, median, or mode.

- **Mean/Median Imputation:** Use for numerical data (mean for normally distributed data, median for skewed data).
- **Mode Imputation:** Use for categorical data.
- **Pros:** Preserves data size.

- **Cons:** Can introduce bias if missing data is not random.

### 3.1.3 Using Algorithms to Predict Missing Values

Use machine learning (e.g., KNN) to predict missing values based on other features.

- **Pros:** More accurate than simple imputation.
- **Cons:** Computationally expensive.

#### Code Example:

```

1 import pandas as pd
2 import numpy as np
3 from sklearn.impute import SimpleImputer
4
5 # Sample dataset with missing values
6 data = {
7     'Age': [25, np.nan, 30, 35, np.nan],
8     'Salary': [50000, 60000, np.nan, 80000, 70000],
9     'Gender': ['Male', 'Female', np.nan, 'Male', 'Female']
10 }
11 df = pd.DataFrame(data)
12
13 # 1. Remove rows with missing values
14 df_dropped = df.dropna()
15 print("After dropping rows with missing values:\n", df_dropped)
16
17 # 2. Impute numerical columns with mean
18 num_imputer = SimpleImputer(strategy='mean')
19 df[['Age', 'Salary']] = num_imputer.fit_transform(df[['Age',
20     'Salary']])
21 print("\nAfter mean imputation for Age and Salary:\n", df)
22
23 # 3. Impute categorical column with mode
24 cat_imputer = SimpleImputer(strategy='most_frequent')
25 df[['Gender']] = cat_imputer.fit_transform(df[['Gender']])
26 print("\nAfter mode imputation for Gender:\n", df)

```

#### Explanation:

- SimpleImputer from Scikit-learn replaces missing values.
- For Age and Salary, we use the mean to fill missing values.
- For Gender, we use the mode (most frequent value).
- Output shows how the dataset changes after each step.

## 3.2 Data Cleaning

Data cleaning involves fixing inconsistencies, duplicates, and errors in the data.

### 3.2.1 Remove Duplicates

Duplicate rows can bias the model by over-representing certain data points. Use Pandas `drop_duplicates()` to remove them.

### 3.2.2 Fix Inconsistent Formats

Standardize text (e.g., convert "USA" and "United States" to one format). Correct typos or case sensitivity (e.g., "male" vs. "Male").

### 3.2.3 Handle Incorrect Data

Replace invalid entries (e.g., negative age) with valid ones or remove them.

#### Code Example:

```
1 # Sample dataset with duplicates and inconsistencies
2 data = {
3     'Name': ['John', 'john', 'Alice', 'Bob', 'Bob'],
4     'Age': [25, 25, 30, -5, 35],
5     'Country': ['USA', 'United States', 'Canada', 'USA', 'USA']
6 }
7 df = pd.DataFrame(data)
8
9 # 1. Remove duplicates
10 df = df.drop_duplicates()
11 print("After removing duplicates:\n", df)
12
13 # 2. Standardize 'Country' column
14 df['Country'] = df['Country'].replace('United States', 'USA')
15 print("\nAfter standardizing Country:\n", df)
16
17 # 3. Fix invalid ages (e.g., negative values)
18 df.loc[df['Age'] < 0, 'Age'] = np.nan # Replace negative ages with NaN
19 df['Age'] = df['Age'].fillna(df['Age'].mean()) # Impute with mean
20 print("\nAfter fixing invalid ages:\n", df)
```

#### Explanation:

- `drop_duplicates()` removes identical rows (e.g., two "Bob" entries).
- `replace()` standardizes "United States" to "USA".
- Negative ages are replaced with NaN, then imputed with the mean.

## 3.3 Encoding Categorical Variables

Machine learning models require numerical inputs, but datasets often contain categorical data (e.g., "Red", "Blue"). Encoding converts these into numbers.

### 3.3.1 Label Encoding

Assigns a unique integer to each category (e.g., "Male" → 0, "Female" → 1).

- **When to Use:** For ordinal data (where order matters, e.g., "Low", "Medium", "High").
- **Cons:** Can mislead models if used on non-ordinal data.

### 3.3.2 One-Hot Encoding

Creates binary columns for each category (e.g., "Color" with "Red", "Blue" becomes "Color\_Red", "Color\_Blue").

- **When to Use:** For nominal data (no order, e.g., colors, countries).

- **Cons:** Increases dataset size with many categories.

### 3.3.3 Target Encoding

Replaces categories with the mean of the target variable for that category.

- **When to Use:** For high-cardinality categorical features (many unique values).

**Code Example:**

```

1 from sklearn.preprocessing import LabelEncoder, OneHotEncoder
2
3 # Sample dataset
4 data = {
5     'Color': ['Red', 'Blue', 'Green', 'Red'],
6     'Size': ['Small', 'Medium', 'Large', 'Medium']
7 }
8 df = pd.DataFrame(data)
9
10 # 1. Label Encoding for 'Size' (ordinal)
11 label_encoder = LabelEncoder()
12 df['Size_Encoded'] = label_encoder.fit_transform(df['Size'])
13 print("After label encoding Size:\n", df)
14
15 # 2. One-Hot Encoding for 'Color' (nominal)
16 one_hot_encoder = OneHotEncoder(sparse_output=False)
17 color_encoded = one_hot_encoder.fit_transform(df[['Color']])
18 color_df = pd.DataFrame(color_encoded,
19                          columns=one_hot_encoder.get_feature_names_out(['Color']))
20 df = pd.concat([df, color_df], axis=1)
21 print("\nAfter one-hot encoding Color:\n", df)

```

**Explanation:**

- `LabelEncoder` assigns numbers to "Small", "Medium", "Large" (e.g., 0, 1, 2).
- `OneHotEncoder` creates binary columns for "Red", "Blue", "Green".
- The resulting `DataFrame` includes both encoded features.

## 3.4 Feature Scaling

Features with different scales (e.g., age: 0-100, salary: 0-100000) can bias models, as larger values may dominate. Scaling standardizes feature ranges.

### 3.4.1 Standardization (Z-Score Normalization)

Scales features to have a mean of 0 and a standard deviation of 1.

- **Formula:**  $z = \frac{x - \text{mean}}{\text{std}}$
- **When to Use:** For algorithms like SVM, logistic regression, or neural networks.

### 3.4.2 Min-Max Scaling

Scales features to a fixed range, typically [0, 1].

- **Formula:**  $x' = \frac{x - \min}{\max - \min}$
- **When to Use:** For algorithms sensitive to feature ranges, like k-nearest neighbors.

### Code Example:

```
1 from sklearn.preprocessing import StandardScaler, MinMaxScaler
2
3 # Sample dataset
4 data = {
5     'Age': [25, 30, 35, 40, 45],
6     'Salary': [50000, 60000, 75000, 80000, 90000]
7 }
8 df = pd.DataFrame(data)
9
10 # 1. Standardization
11 std_scaler = StandardScaler()
12 df[['Age_Std', 'Salary_Std']] = std_scaler.fit_transform(df[['Age',
13     'Salary']])
14 print("After standardization:\n", df)
15
16 # 2. Min-Max Scaling
17 minmax_scaler = MinMaxScaler()
18 df[['Age_MinMax', 'Salary_MinMax']] =
19     minmax_scaler.fit_transform(df[['Age', 'Salary']])
20 print("\nAfter Min-Max scaling:\n", df)
```

### Explanation:

- `StandardScaler` transforms Age and Salary to have mean=0 and std=1.
- `MinMaxScaler` scales them to [0, 1].
- The output shows both scaled versions for comparison.

## 3.5 Handling Outliers

Outliers are extreme values that can skew model performance. Common methods to handle them:

### 3.5.1 Remove Outliers

Use statistical methods like the Interquartile Range (IQR) to detect and remove outliers.

- **IQR Method:** Outliers are values below  $Q1 - 1.5 \cdot \text{IQR}$  or above  $Q3 + 1.5 \cdot \text{IQR}$ .

### 3.5.2 Cap Outliers

Replace outliers with the nearest acceptable value (e.g., cap at the 95th percentile).

### 3.5.3 Transform Features

Apply transformations like log or square root to reduce the impact of outliers.

### Code Example:

```
1 # Sample dataset with outliers
2 data = {
3     'Salary': [50000, 60000, 75000, 1000000, 80000]
4 }
5 df = pd.DataFrame(data)
6
7 # 1. Detect and remove outliers using IQR
```

```

8 Q1 = df['Salary'].quantile(0.25)
9 Q3 = df['Salary'].quantile(0.75)
10 IQR = Q3 - Q1
11 lower_bound = Q1 - 1.5 * IQR
12 upper_bound = Q3 + 1.5 * IQR
13 df_no_outliers = df[(df['Salary'] >= lower_bound) & (df['Salary'] <=
    upper_bound)]
14 print("After removing outliers:\n", df_no_outliers)
15
16 # 2. Cap outliers
17 df['Salary_Capped'] = df['Salary'].clip(lower=lower_bound,
    upper=upper_bound)
18 print("\nAfter capping outliers:\n", df)

```

### Explanation:

- The IQR method identifies 1000000 as an outlier and removes it.
- clip() caps extreme values to the upper/lower bounds.

## 3.6 Feature Selection

Feature selection removes irrelevant or redundant features to improve model performance and reduce training time.

### 3.6.1 Filter Methods

Select features based on statistical measures (e.g., correlation with the target).

- **Example:** Remove features with low variance or high correlation.

### 3.6.2 Wrapper Methods

Evaluate subsets of features using a model (e.g., recursive feature elimination).

### 3.6.3 Embedded Methods

Use algorithms that inherently perform feature selection (e.g., Lasso regression).

### Code Example:

```

1 from sklearn.feature_selection import VarianceThreshold
2
3 # Sample dataset
4 data = {
5     'Feature1': [1, 1, 1, 1], # Low variance
6     'Feature2': [2, 4, 6, 8], # High variance
7     'Feature3': [1, 2, 3, 4]
8 }
9 df = pd.DataFrame(data)
10
11 # Remove low-variance features
12 selector = VarianceThreshold(threshold=0.1)
13 df_selected = selector.fit_transform(df)
14 print("Selected features (after removing low variance):\n",
    pd.DataFrame(df_selected,
        columns=df.columns[selector.get_support()]))

```

### Explanation:

- `VarianceThreshold` removes `Feature1` because it has low variance (all values are 1).

## 3.7 Feature Engineering

Feature engineering creates new features or transforms existing ones to improve model performance.

### 3.7.1 Creating New Features

Combine features (e.g., create "Age\_Salary\_Ratio" from `Age` and `Salary`). Extract features (e.g., extract "Day" from a date column).

### 3.7.2 Transforming Features

Apply mathematical transformations (e.g., log, square). Bin continuous features into categories (e.g., age groups).

### Code Example:

```
1 # Sample dataset
2 data = {
3     'Age': [25, 30, 35, 40],
4     'Salary': [50000, 60000, 75000, 80000],
5     'Date': ['2023-01-01', '2023-02-01', '2023-03-01', '2023-04-01']
6 }
7 df = pd.DataFrame(data)
8
9 # 1. Create new feature: Age_Salary_Ratio
10 df['Age_Salary_Ratio'] = df['Age'] / df['Salary']
11
12 # 2. Extract month from Date
13 df['Date'] = pd.to_datetime(df['Date'])
14 df['Month'] = df['Date'].dt.month
15
16 # 3. Bin Age into categories
17 df['Age_Group'] = pd.cut(df['Age'], bins=[0, 30, 40, 100],
18                           labels=['Young', 'Middle', 'Senior'])
19 print("After feature engineering:\n", df)
```

### Explanation:

- `Age_Salary_Ratio` is a new feature combining `Age` and `Salary`.
- `dt.month` extracts the month from the `Date` column.
- `pd.cut` bins `Age` into categories like "Young" and "Middle".

## 3.8 Data Splitting

Before training a model, split the dataset into training, validation, and test sets to evaluate performance.

- **Training Set:** Used to train the model (e.g., 70% of data).
- **Validation Set:** Used to tune hyperparameters (e.g., 15% of data).
- **Test Set:** Used to evaluate final performance (e.g., 15% of data).



### Code Example:

```
1 from sklearn.model_selection import train_test_split
2
3 # Sample dataset
4 data = {
5     'Feature1': [1, 2, 3, 4, 5, 6, 7, 8, 9, 10],
6     'Target': [0, 1, 0, 1, 0, 1, 0, 1, 0, 1]
7 }
8 df = pd.DataFrame(data)
9
10 # Split data
11 X = df[['Feature1']]
12 y = df['Target']
13 X_train, X_temp, y_train, y_temp = train_test_split(X, y,
14     test_size=0.3, random_state=42)
15 X_val, X_test, y_val, y_test = train_test_split(X_temp, y_temp,
16     test_size=0.5, random_state=42)
17
18 print("Training set size:", len(X_train))
19 print("Validation set size:", len(X_val))
20 print("Test set size:", len(X_test))
```

### Explanation:

- `train_test_split` splits the data into 70% training and 30% temporary sets.
- The temporary set is further split into 15% validation and 15% test sets.

## 4 Putting It All Together: A Complete Example

Heres a complete preprocessing pipeline combining all techniques.

### Code Example:

```
1 import pandas as pd
2 import numpy as np
3 from sklearn.impute import SimpleImputer
4 from sklearn.preprocessing import StandardScaler, OneHotEncoder
5 from sklearn.model_selection import train_test_split
6
7 # Sample dataset
8 data = {
9     'Age': [25, np.nan, 30, 35, 1000, 40],
10    'Salary': [50000, 60000, np.nan, 80000, 90000, 85000],
11    'Gender': ['Male', 'Female', np.nan, 'Male', 'Female', 'Male'],
12    'Country': ['USA', 'Canada', 'USA', 'United States', 'Canada',
13        'USA'],
14    'Target': [0, 1, 0, 1, 0, 1]
15 }
16 df = pd.DataFrame(data)
17
18 # 1. Handle missing values
19 num_imputer = SimpleImputer(strategy='mean')
20 df[['Age', 'Salary']] = num_imputer.fit_transform(df[['Age',
21     'Salary']])
22 cat_imputer = SimpleImputer(strategy='most_frequent')
23 df[['Gender']] = cat_imputer.fit_transform(df[['Gender']])
```

```

23 # 2. Clean data (standardize Country, handle outliers)
24 df['Country'] = df['Country'].replace('United States', 'USA')
25 Q1 = df['Age'].quantile(0.25)
26 Q3 = df['Age'].quantile(0.75)
27 IQR = Q3 - Q1
28 df['Age'] = df['Age'].clip(lower=Q1 - 1.5 * IQR, upper=Q3 + 1.5 * IQR)
29
30 # 3. Encode categorical variables
31 df = pd.get_dummies(df, columns=['Gender', 'Country'], drop_first=True)
32
33 # 4. Scale features
34 scaler = StandardScaler()
35 df[['Age', 'Salary']] = scaler.fit_transform(df[['Age', 'Salary']])
36
37 # 5. Split data
38 X = df.drop('Target', axis=1)
39 y = df['Target']
40 X_train, X_test, y_train, y_test = train_test_split(X, y,
41                                                    test_size=0.2, random_state=42)
42
43 print("Preprocessed training data:\n", X_train)
44 print("\nPreprocessed test data:\n", X_test)

```

### Explanation:

- The pipeline handles missing values, cleans data, encodes categories, scales features, and splits the data.
- The final output shows the preprocessed training and test sets, ready for model training.

## 5 Best Practices and Tips

### 1. Understand Your Data:

- Explore the dataset using `df.describe()`, `df.info()`, and visualizations to identify issues.
- Check for missing values, outliers, and data types.

### 2. Avoid Data Leakage:

- Apply preprocessing (e.g., scaling, imputation) on the training set first, then apply the same transformation to the test set. Never fit preprocessors on the test set.

### 3. Choose Appropriate Techniques:

- Use mean imputation for normally distributed data, median for skewed data.
- Use one-hot encoding for nominal data, label encoding for ordinal data.

### 4. Automate Preprocessing:

- Use Scikit-learn's `Pipeline` to combine preprocessing steps and ensure consistency.

### 5. Document Your Steps:

- Keep track of preprocessing decisions to reproduce results or debug issues.

### 6. Handle Imbalanced Data (if applicable):

- Use techniques like oversampling (SMOTE) or undersampling for imbalanced datasets.

### Code Example (Pipeline):

```
1 from sklearn.pipeline import Pipeline
2 from sklearn.compose import ColumnTransformer
3
4 # Define numeric and categorical columns
5 numeric_features = ['Age', 'Salary']
6 categorical_features = ['Gender', 'Country']
7
8 # Create preprocessor
9 preprocessor = ColumnTransformer(
10     transformers=[
11         ('num', Pipeline([
12             ('imputer', SimpleImputer(strategy='mean')),
13             ('scaler', StandardScaler()))
14         ]), numeric_features),
15         ('cat', Pipeline([
16             ('imputer', SimpleImputer(strategy='most_frequent')),
17             ('onehot', OneHotEncoder(drop='first'))
18         ]), categorical_features)
19     ])
20
21 # Apply preprocessor
22 X_transformed = preprocessor.fit_transform(df.drop('Target', axis=1))
23 print("Preprocessed data using pipeline:\n",
      pd.DataFrame(X_transformed))
```

### Explanation:

- `ColumnTransformer` applies different preprocessing steps to numeric and categorical columns.
- `Pipeline` ensures consistent application of imputation and scaling.

This guide covers all essential preprocessing techniques with detailed explanations and code. By following these steps, you can prepare any dataset for machine learning. Save this guide, and you'll have a complete reference for preprocessing!